

Monte Carlo Coulomb Collisions for Guiding-Centre Particles in Boozer Coordinates: Implementation Notes

firm3d — ep_thermal_collisions

August 5, 2026

Contents

1	Overview	2
2	Guiding-Centre Equations in Boozer Coordinates	2
3	Background Species and Profiles	3
4	Collision Coefficients	3
4.1	Debye Length	3
4.2	Coulomb Logarithm	3
4.3	Chandrasekhar G Function	4
4.4	Collision Frequency Prefactor	5
4.5	Pitch-Angle Scattering Frequency	5
4.6	Parallel Diffusion Coefficient	5
4.7	Deterministic Drag in v	5
4.8	Summation over Species	6
5	Stochastic Differential Equations	7
5.1	Full SDE System	7
5.2	Milstein Scheme	7
6	Adaptive Integration Strategy	8
6.1	Dormand–Prince Solver	8
6.2	Minimum Step Size	8
6.3	Operator Splitting	8
6.4	Sub-Cycling of the Collision Kick	9
7	Output Format	9
8	Python Interface	9
9	Testing and Validation	10
9.1	Analytic Coefficient Unit Tests	11
9.2	Local Velocity-Space Relaxation Tests	11
9.3	End-to-End Equilibration Tests (Production Tracer)	11

9.4	Slow Physics Suite (MPI, HPC)	12
9.5	Cross-Code Validation Against ASCOT5: Uniform Plasma	13
10	Validity and Limitations	13

1 Overview

These notes describe the algorithm used in `trace_particles_boozer_with_collisions` to include stochastic Coulomb collisions for an energetic particle (EP) tracing through one or more Maxwellian background species.

The approach couples the deterministic guiding-centre (GC) orbit equations in Boozer coordinates with a Milstein stochastic term for the velocity-space diffusion. The integrated state is the ordinary GC state $(s, \theta, \zeta, v_{\parallel})$ at fixed magnetic moment μ ; the collision operator acts on the velocity-space pair (v, ξ) , where $v = |\mathbf{v}|$ is the total speed and $\xi = v_{\parallel}/v$ the pitch, as a kick applied between orbit steps (Section 6.3). Because μ is constant across an orbit step it is a parameter of the orbit equations rather than a state variable, which is what allows the vacuum, $K = 0$ and full GC equations all to be driven by the same collision operator.

Primary references:

- Hirvijoki et al. [1] — Monte Carlo implementation of the guiding-centre Fokker–Planck equation in (v, ξ) coordinates; Eqs. (32)–(33) of that paper define the SDE system used here.
- Hirvijoki et al. [2] — ASCOT4/5 code description; source file `mccc_coefs.h` defines the Coulomb logarithm convention adopted here.
- Boozer & Kuo-Petravic [4] — pitch-angle representation in Boozer coordinates.

2 Guiding-Centre Equations in Boozer Coordinates

In Boozer coordinates (s, θ, ζ) , where s is the normalised toroidal flux, θ the poloidal angle, and ζ the toroidal angle, the vacuum GC equations are

$$\dot{s} = -\frac{m v^2 (1 + \xi^2)}{2 q \psi_0} \frac{\partial B}{\partial \theta} \frac{1}{B}, \quad (1)$$

$$\dot{\theta} = \frac{m v^2 (1 + \xi^2)}{2 q \psi_0} \frac{\partial B}{\partial s} \frac{1}{B} + \frac{\iota \xi v B}{G}, \quad (2)$$

$$\dot{\zeta} = \frac{\xi v B}{G}, \quad (3)$$

$$\dot{\xi} = -\frac{(\iota \partial_{\theta} B + \partial_{\zeta} B) v (1 - \xi^2)}{2G} \quad [\text{mirror force}], \quad (4)$$

where ψ_0 is the reference poloidal flux, $G(s)$ is the covariant toroidal field component, and $\iota(s)$ is the rotational transform. The magnetic moment $\mu = v^2(1 - \xi^2)/(2B)$ is held fixed through an orbit step and updated by the collision kick. The vacuum equations are shown here for

concreteness; the $K = 0$ and full GC equations of `trace_particles_boozer` are supported identically and are selected from the field's `field_type`.

With collisions the equations for v and ξ acquire deterministic drift terms (Section 4) and stochastic noise terms (Section 5.2).

3 Background Species and Profiles

Each background species b is described by a `ThermalBackground` object specifying:

- $n_b(s)$ — number density [m^{-3}], callable of the flux label $s \in [0, 1]$;
- $T_b(s)$ — temperature [J], callable (use `T_keV × 1e3 × e`);
- m_b — species mass [kg];
- q_b — species charge [C] (signed).

Internally the profiles are pre-evaluated on a uniform s -grid of 512 points and linearly interpolated in C++.

4 Collision Coefficients

Notation

The symbols used here follow the ASCOT5 source code (`mccc_coefs.h`, [3]); Table 1 shows the correspondence with Hirvijoki et al. [1], whose Eqs. (32)–(33) define the SDE structure.

This paper	ASCOT5 <code>mccc_coefs.h</code>	Hirvijoki et al. [1]
$D_{\parallel}$	<code>Dpara</code>	$D_k/2$
ν_D	<code>nu</code> = $2D_{\perp}/v^2$	$2D_{\perp}/v^2$
Q	<code>Q</code>	(Einstein-relation drag in K)
K	<code>K</code> = $Q + D'_{\parallel} + 2D_{\parallel}/v$	K (drift in Eq. 32)

Table 1: Notation correspondence. The noise amplitude for v is $g_v = \sqrt{2D_{\parallel}}$ (this paper) = $\sqrt{D_k}$ (Hirvijoki 2013).

4.1 Debye Length

The total Debye screening length from all background species is [5, 6]

$$\frac{1}{\lambda_D^2} = \sum_b \frac{n_b q_b^2}{\varepsilon_0 T_b}. \quad (5)$$

4.2 Coulomb Logarithm

The Coulomb logarithm is defined as the ratio of maximum to minimum impact parameters [5, 6, 7]:

$$\ln \Lambda_{ab} = \ln \left(\frac{\lambda_D}{b_{\min}} \right), \quad b_{\min} = \max(b_{\text{cl}}, b_{\text{qm}}). \quad (6)$$

The maximum impact parameter is the Debye length (Eq. (5)). The minimum impact parameter is the larger of the classical 90°-deflection distance and the quantum-mechanical de Broglie wavelength of the reduced mass:

$$b_{\text{cl}} = \frac{|q_a q_b|}{4\pi\epsilon_0 m_r \bar{v}^2}, \quad (7)$$

$$b_{\text{qm}} = \frac{\hbar}{2 m_r \sqrt{\bar{v}^2}}, \quad (8)$$

where $m_r = m_a m_b / (m_a + m_b)$ is the reduced mass.

The effective velocity scale $\bar{v}^2$ is taken as

$$\bar{v}^2 = v^2 + v_{\text{th},b}^2, \quad v_{\text{th},b} = \sqrt{2T_b/m_b}, \quad (9)$$

following the ASCOT5 source convention [3]. This form is a numerical interpolation device, not derived from first principles: Ichimaru [7] derives b_{cl} and b_{qm} only in the fast-particle limit $v \gg v_{\text{th},b}$ (Section 4.3B of that reference), where $\bar{v} \approx v$. The $\bar{v}^2$ form extends the formula to the slow-particle regime without a discontinuous switch: fast EP against ions ($v \gg v_{\text{th},b}$, $\bar{v} \approx v$) and slow EP against electrons ($v \ll v_{\text{th},e}$, $\bar{v} \approx v_{\text{th},e}$).

The quantum branch $b_{\text{qm}} > b_{\text{cl}}$ applies when

$$\frac{2\pi\epsilon_0 \hbar \sqrt{\bar{v}^2}}{|q_a q_b|} > 1, \quad (10)$$

i.e. when the de Broglie wavelength of the reduced mass particle exceeds the classical distance of closest approach. For 3.52 MeV α -particles scattering off 10 keV electrons one finds $b_{\text{qm}} \approx 6.9 b_{\text{cl}}$, so the quantum branch is active throughout alpha slowing-down against electrons. If $\ln \Lambda \leq 0$ (Debye length below the minimum impact parameter, e.g. $T_b \rightarrow 0$ at finite density) the binary-collision model is undefined and an exception is raised; ASCOT5 reaches the equivalent outcome by aborting markers through its input-evaluation errors, and its input profiles keep the temperature finite wherever the density is nonzero. A warning is issued for $0 < \ln \Lambda < 2$, where the perturbative expansion underlying the Fokker–Planck operator begins to lose validity.

Both checks are applied to the profiles once, before tracing starts, rather than from inside the right-hand side: the latter is evaluated roughly seven times per accepted step per particle, so a per-evaluation warning would emit millions of duplicate lines, and a per-evaluation exception would fire on whichever MPI rank owns the offending particle, hanging the remaining ranks in the collective that gathers results. Since $\ln \Lambda$ increases monotonically with v through $\bar{v}^2$, evaluating at $v \rightarrow 0$ bounds it from below over all particle speeds, making the up-front check conservative but never optimistic.

4.3 Chandrasekhar G Function

The Chandrasekhar function [8] is

$$G(x) = \frac{\text{erf}(x) - \frac{2x}{\sqrt{\pi}} e^{-x^2}}{2x^2}, \quad x = v/v_{\text{th},b}. \quad (11)$$

Its derivative satisfies the recurrence

$$G'(x) = \frac{2}{\sqrt{\pi}} e^{-x^2} - \frac{2G(x)}{x}. \quad (12)$$

Key limiting behaviours:

$$\begin{aligned} x \rightarrow 0 : \quad G &\rightarrow 0, \quad G' \rightarrow 2/(3\sqrt{\pi}) \approx 0.376; \\ x \rightarrow \infty : \quad G &\rightarrow 1/(2x^2), \quad G' \rightarrow 0. \end{aligned}$$

4.4 Collision Frequency Prefactor

Define the Rosenbluth prefactor for species b [9, 10]:

$$\Gamma_b = \frac{n_b q_a^2 q_b^2 \ln \Lambda_{ab}}{4\pi\epsilon_0^2 m_a^2}. \quad (13)$$

4.5 Pitch-Angle Scattering Frequency

$$\nu_D^{(b)} = \frac{\Gamma_b [\operatorname{erf}(x) - G(x)]}{v^3} \equiv \frac{2D_{\perp}^{(b)}}{v^2}. \quad (14)$$

The second form uses the perpendicular diffusion coefficient $D_{\perp}^{(b)} = \Gamma_b [\operatorname{erf}(x) - G(x)]/(2v)$ and matches ASCOT5's `mccc_coefs_nu` and the $2D_{\perp}/v^2$ notation of Hirvijoki et al. [1, 2].

4.6 Parallel Diffusion Coefficient

$$D_{\parallel}^{(b)} = \frac{\Gamma_b G(x)}{v}, \quad \frac{\partial D_{\parallel}^{(b)}}{\partial v} = \frac{\Gamma_b}{v} \left(\frac{G'(x)}{v_{\text{th},b}} - \frac{G(x)}{v} \right). \quad (15)$$

In Hirvijoki et al. [1] the parallel diffusion coefficient is denoted D_k , with $D_{\parallel} = D_k/2$; the noise amplitude for the speed SDE is $\sqrt{D_k} = \sqrt{2D_{\parallel}}$ in both conventions. ASCOT5 names this coefficient `Dpara`. The v -derivative is required by the Milstein correction term (Section 5.2). Note that it differentiates Γ_b as though $\ln \Lambda_{ab}$ were constant, even though the $\ln \Lambda_{ab}$ actually evaluated depends on v through $\bar{v}^2$; this follows ASCOT5 and has a small but systematic consequence for the equilibrium, quantified in Section 10.

4.7 Deterministic Drag in v

The drag coefficient (ASCOT5: `mccc_coefs_Q` [3]) is fixed by the Einstein relation:

$$Q^{(b)} = -\frac{m_a v}{T_b} D_{\parallel}^{(b)} = -\Gamma_b \frac{m_a G(x)}{T_b}. \quad (16)$$

The total deterministic drift in speed is

$$K^{(b)} = Q^{(b)} + \frac{\partial D_{\parallel}^{(b)}}{\partial v} + \frac{2 D_{\parallel}^{(b)}}{v} \quad (17)$$

(ASCOT5's `mccc_coefs_K` [3]; the drift coefficient K of Eq. (32) in Hirvijoki et al. [1]).

Itô transformation leading to Eq. (17)

Equation (17) can be obtained two independent ways, which cross-check each other.

(i) **Itô's formula on $v = |\mathbf{v}|$.** In three-dimensional velocity space the collisional motion of the test particle is

$$d\mathbf{v} = F_{\text{dyn}} \hat{\mathbf{v}} dt + \boldsymbol{\sigma} \cdot d\mathbf{W}, \quad \boldsymbol{\sigma} \boldsymbol{\sigma}^\top = 2D_{\parallel} \hat{\mathbf{v}} \hat{\mathbf{v}} + 2D_{\perp} (\mathbf{I} - \hat{\mathbf{v}} \hat{\mathbf{v}}), \quad (18)$$

where F_{dyn} is the dynamical friction [8, 10, 11]. The speed $v = |\mathbf{v}|$ is a nonlinear function of $\mathbf{v}$, so Itô's formula applies with $\nabla_{\mathbf{v}} v = \hat{\mathbf{v}}$ and Hessian $\nabla_{\mathbf{v}} \nabla_{\mathbf{v}} v = (\mathbf{I} - \hat{\mathbf{v}} \hat{\mathbf{v}})/v$:

$$dv = \left(F_{\text{dyn}} + \frac{2D_{\perp}}{v} \right) dt + \sqrt{2D_{\parallel}} dW_v. \quad (19)$$

The $2D_{\perp}/v$ drift is the Bessel-process effect: kicks transverse to $\mathbf{v}$ do not change v at first order, but they lengthen the vector at second order, $\Delta v = |\mathbf{v} + \Delta \mathbf{v}_{\perp}| - v \approx |\Delta \mathbf{v}_{\perp}|^2/2v > 0$.

(ii) **Fokker–Planck matching (detailed balance).** The Itô SDE $dv = K dt + \sqrt{2D_{\parallel}} dW_v$ evolves the speed density $P(v) \propto v^2 f(v)$ according to $\partial_t P = -\partial_v(KP) + \partial_v^2(D_{\parallel}P)$. Zero flux on the equilibrium $P_{\text{eq}} \propto v^2 e^{-m_a v^2/2T_b}$ requires $KP_{\text{eq}} = \partial_v(D_{\parallel}P_{\text{eq}})$, i.e.

$$K = -\frac{m_a v}{T_b} D_{\parallel} + \frac{\partial D_{\parallel}}{\partial v} + \frac{2D_{\parallel}}{v}, \quad (20)$$

which is Eq. (17) with Q as in Eq. (16). Here $\partial_v D_{\parallel}$ is the multiplicative-noise (Itô) drift that appears when the diffusion term is commuted into advective form, and $2D_{\parallel}/v$ comes from the Jacobian v^2 of spherical velocity-space coordinates.

Consistency of (i) and (ii). For Maxwellian field particles the dynamical friction obeys the identity

$$F_{\text{dyn}}^{(b)} = -\left(1 + \frac{m_a}{m_b}\right) \frac{2x^2 G(x) \Gamma_b}{v^2} = Q^{(b)} + \frac{\partial D_{\parallel}^{(b)}}{\partial v} + \frac{2(D_{\parallel}^{(b)} - D_{\perp}^{(b)})}{v}, \quad (21)$$

so that $F_{\text{dyn}} + 2D_{\perp}/v = Q + \partial_v D_{\parallel} + 2D_{\parallel}/v = K$ exactly. (F_{dyn} is ASCOT5's `mccc_coefs_F`, which is not used in the guiding-centre scheme.) In the fast-particle limit $x \gg 1$ this identity encodes the classic cancellation

$$F_{\text{dyn}} + \frac{2D_{\perp}}{v} \longrightarrow -\left(1 + \frac{m_a}{m_b}\right) \frac{\Gamma_b}{v^2} + \frac{\Gamma_b}{v^2} = -\frac{m_a}{m_b} \frac{\Gamma_b}{v^2} :$$

most of the dynamical friction is compensated by transverse diffusion, and the net slowing-down rate is set by the mass ratio. Note in particular that Q is *not* the dynamical friction — it is the Einstein-relation drag $-(m_a v/T_b) D_{\parallel}$.

4.8 Summation over Species

All coefficients are additive over background species:

$$\nu_D = \sum_b \nu_D^{(b)}, \quad D_{\parallel} = \sum_b D_{\parallel}^{(b)}, \quad K = \sum_b K^{(b)}. \quad (22)$$

The Debye length used in Eq. (6) is computed from all species simultaneously (Eq. (5)), prior to the per-species loop.

5 Stochastic Differential Equations

5.1 Full SDE System

The system being solved is the following, written in the natural velocity-space variables [1]. Section 6.3 describes how it is split between the orbit stepper and the collision kick.

$$ds = \dot{s}_{\text{orb}} dt, \quad (23)$$

$$d\theta = \dot{\theta}_{\text{orb}} dt, \quad (24)$$

$$d\zeta = \dot{\zeta}_{\text{orb}} dt, \quad (25)$$

$$dv = K(v, s) dt + \sqrt{2D_{\parallel}} dW_v, \quad (26)$$

$$d\xi = \dot{\xi}_{\text{mirror}} dt - \xi \nu_D dt + \sqrt{\nu_D(1 - \xi^2)} dW_{\xi}, \quad (27)$$

where dW_v and dW_{ξ} are independent Wiener increments with $\mathbb{E}[dW^2] = dt$. Equations (26)–(27) correspond to Eqs. (32)–(33) of Hirvijoki et al. [1], with the identification $D_{\parallel} \leftrightarrow D_k/2$ and $\nu_D \leftrightarrow 2D_{\perp}/v^2$ in their notation. The noise amplitude for ξ ensures that the marginal distribution of ξ converges to a uniform distribution on $[-1, 1]$ in the isotropisation limit $\nu_D t \gg 1$.

5.2 Milstein Scheme

After each accepted Dormand–Prince (DP) step of physical size h , the stochastic increments $\Delta W_v \sim \mathcal{N}(0, h)$ and $\Delta W_{\xi} \sim \mathcal{N}(0, h)$ are generated and applied via the Milstein method [12]:

$$v \leftarrow v + K h + g_v \Delta W_v + \frac{1}{2} \frac{\partial g_v}{\partial v} (\Delta W_v^2 - h), \quad (28)$$

$$\xi \leftarrow \xi - \xi \nu_D h + g_{\xi} \Delta W_{\xi} + \frac{1}{2} \frac{\partial g_{\xi}}{\partial \xi} (\Delta W_{\xi}^2 - h), \quad (29)$$

The deterministic drift terms Kh and $-\xi \nu_D h$ are applied here, by explicit Euler, rather than inside the orbit right-hand side; this matches ASCOT5’s `mccc_gc_milstein.c` and is what keeps μ constant across an orbit step. where the noise amplitudes and their derivatives are

$$g_v = \sqrt{2D_{\parallel}}, \quad \frac{\partial g_v}{\partial v} = \frac{\partial D_{\parallel}/\partial v}{\sqrt{2D_{\parallel}}}, \quad (30)$$

$$g_{\xi} = \sqrt{\nu_D(1 - \xi^2)}, \quad \frac{\partial g_{\xi}}{\partial \xi} = \frac{-\xi \nu_D}{\sqrt{\nu_D(1 - \xi^2)}}. \quad (31)$$

In terms of the coefficient structs, the Milstein correction for ξ simplifies to

$$\delta \xi_{\text{Milstein}} = -\frac{1}{2} \xi \nu_D (\Delta W_{\xi}^2 - h). \quad (32)$$

Boundary conditions follow ASCOT5 (`mccc_gc_milstein.c`, [3]). The speed reflects off a thermal cutoff,

$$v < v_{\text{cut}} \Rightarrow v \leftarrow 2v_{\text{cut}} - v, \quad v_{\text{cut}} = 0.1 \sqrt{T_{\text{min}}/m_a}, \quad (33)$$

where T_{min} is the coldest active background species (ASCOT5’s `MCCC_CUTOFF` uses its first background species), so that the collision coefficients, which diverge as $v \rightarrow 0$, are never evaluated deep below the thermal speed. The pitch mirrors at the boundary: $|\xi| > 1 \Rightarrow \xi \leftarrow \text{sign}(\xi) (2 - |\xi|)$; a hard clamp would pile up probability at exactly $\xi = \pm 1$.

Order of convergence

Only the diagonal Milstein corrections are retained above. In more than one dimension these alone give order-1 strong convergence just when the diffusion matrix is commutative, i.e. $L^{j_1} g^{k,j_2} = L^{j_2} g^{k,j_1}$ with $L^j = \sum_k g^{k,j} \partial_k$ [12]. The diffusion here is diagonal — dW_v drives v and dW_ξ drives ξ — so with $k = \xi$,

$$L^v g_\xi = g_v \frac{\partial g_\xi}{\partial v} \neq 0, \quad L^\xi g_v = g_\xi \frac{\partial g_v}{\partial \xi} = 0, \quad (34)$$

the first being nonzero because ν_D depends on v , and the second vanishing because $D_\parallel$ does not depend on ξ . The condition fails, so the double stochastic integrals (Lévy areas) $I_{(v,\xi)}$ and $I_{(\xi,v)}$ do not cancel and cannot be recovered from $\Delta W_v, \Delta W_\xi$ alone. Dropping them — as ASCOT5 also does — reduces *strong* convergence to order 1/2 in that coupling. *Weak* convergence, which governs the moments and distribution functions a Monte Carlo transport calculation actually reports, remains order 1 and is unaffected.

6 Adaptive Integration Strategy

6.1 Dormand–Prince Solver

The orbit equations — the four components $(\dot{s}, \dot{\theta}, \dot{\zeta}, \dot{v}_\parallel)$ at fixed μ , with no collision terms — are integrated with the classic embedded Dormand–Prince (FSAL) method [13]. Step-size control uses mixed absolute/relative tolerances:

$$\text{err}_i = |y_i^{(5)} - y_i^{(4)}| \leq \delta_{\text{abs}} + \delta_{\text{rel}} |y_i|. \quad (35)$$

Default tolerances: $\delta_{\text{abs}} = \delta_{\text{rel}} = 10^{-9}$.

6.2 Minimum Step Size

The parameter `DP_hmin` (default 0.0) sets a lower bound on the physical step size h in seconds. A value of 10^{-10} – 10^{-9} s prevents the solver from stalling at bounce points while keeping steps small enough that the orbital dynamics remain accurate. Note this bounds the *orbit* stepper only: the collision terms are no longer part of the right-hand side, so the stiffness they introduce as $v \rightarrow 0$ is handled by the sub-cycling of Section 6.4, not by this floor.

6.3 Operator Splitting

The system is solved via Lie–Trotter operator splitting:

1. **Orbit half:** advance $(s, \theta, \zeta, v_\parallel)$ by one adaptive DP step of physical time h at fixed μ , using the orbit drift only. No collision term appears in this right-hand side.
2. **Collision half:** at the accepted state, convert $(v_\parallel, \mu) \rightarrow (v, \xi)$ using $|B|$ at the current position, apply the drift and Milstein noise of Section 5.2 over the window h , and convert back, which yields the updated μ .

Placing the whole collision operator in the second half is what makes μ exactly conserved by the first, so μ can be a parameter of the orbit equations rather than a state variable. Any static-field GC right-hand side is then usable unchanged.

The trajectory is saved *before* the collision kick at each t_{save} checkpoint, so saved states correspond to the deterministic prediction (consistent with the standard output convention).

6.4 Sub-Cycling of the Collision Kick

The orbit stepper sizes h from orbit dynamics alone, so the collision terms have no influence on it. Applying them as a single explicit-Euler kick over that whole step is inaccurate whenever the collision rates are fast compared with h — the thermal regime, where ν_D and K diverge as $v \rightarrow 0$. Left uncorrected the equilibrium tracks the orbit tolerance instead of the physics: the Maxwellian-equilibration test returned $\langle E \rangle / T_b = 7.8, 3.1, 3.0, 1.3$ at $\delta = 10^{-6}, 10^{-8}, 10^{-10}, 10^{-12}$, against an expected $3/2$.

The kick is therefore sub-divided into n sub-steps, with

$$n = \max \left(\frac{\nu_D h}{\varepsilon}, \frac{|K|h}{\varepsilon v}, \left[\frac{\sqrt{2D_{\parallel}h}}{\varepsilon v} \right]^2 \right), \quad \varepsilon = 0.05, \quad (36)$$

capped at 10^4 . The two drift terms scale as h and so fall linearly in n ; the diffusive excursion scales as $\sqrt{h}$ and so falls only as $\sqrt{n}$, which is why it enters squared. The coefficients are re-evaluated and n re-derived from the remaining time after each sub-step, since v changes within the kick. Splitting is exact for the driving noise: a Wiener increment over h is the sum of n independent increments over h/n .

Sub-cycling is cheap because the particle position is frozen during a kick, so a sub-step costs a coefficient evaluation and no field evaluation. For 3.52 MeV alphas $n = 1$ throughout; only thermal-speed particles sub-cycle. ASCOT5 reaches the same place from the other direction, with error estimates that shorten its collision step (`mccc_gc_milstein.c`: `kappa_k` for drift, `kappa_d0/kappa_d1` for diffusion).

7 Output Format

Each particle trace is a NumPy array of shape (ntimesteps, 6) with columns

$$[t, s, \theta, \zeta, v_{\parallel}, v].$$

From these the kinetic energy $E = \frac{1}{2}mv^2$ and the magnetic moment $\mu = (v^2 - v_{\parallel}^2)/(2B)$ can be reconstructed at each saved point. When `forget_exact_path=True` only the first and last rows are retained.

Each particle uses RNG seed `rng_seed+i`, making multi-particle MPI runs fully reproducible.

8 Python Interface

```
from firm3d.field.collisions import ThermalBackground, \
    trace_particles_boozer_with_collisions
from firm3d.util.constants import (
    PROTON_MASS, ELECTRON_MASS, ELEMENTARY_CHARGE, ONE_EV,
    ALPHA_PARTICLE_MASS, ALPHA_PARTICLE_CHARGE,
    FUSION_ALPHA_PARTICLE_ENERGY,
)

# Typical DT-fusion plasma backgrounds
n0 = 1e20          # m^-3
```

```

deuterium = ThermalBackground(
    n_profile=lambda s: n0 * (1 - s),
    T_profile=lambda s: 10e3 * ONE_EV * (1 - s),
    mass=2 * PROTON_MASS,
    charge=ELEMENTARY_CHARGE,
)
tritium = ThermalBackground(
    n_profile=lambda s: n0 * (1 - s),
    T_profile=lambda s: 10e3 * ONE_EV * (1 - s),
    mass=3 * PROTON_MASS,
    charge=ELEMENTARY_CHARGE,
)
electrons = ThermalBackground(
    n_profile=lambda s: 2 * n0 * (1 - s),    # quasi-neutral
    T_profile=lambda s: 10e3 * ONE_EV * (1 - s),
    mass=ELECTRON_MASS,
    charge=-ELEMENTARY_CHARGE,
)

res_tys, res_hits = trace_particles_boozer_with_collisions(
    field,
    stz_inits,                                # (N, 3) array
    parallel_speeds,                          # (N,) array, m/s
    backgrounds=[deuterium, tritium, electrons],
    tmax=1e-1,                                # 100 ms
    mass=ALPHA_PARTICLE_MASS,
    charge=ALPHA_PARTICLE_CHARGE,
    Ekin=FUSION_ALPHA_PARTICLE_ENERGY,       # 3.52 MeV
    tol=1e-9,
    dt_save=1e-4,                             # save every 0.1 ms
    DP_hmin=1e-10,                            # s
    rng_seed=42,
)

```

9 Testing and Validation

The implementation is exercised at four levels of increasing scope: analytic unit tests of the coefficient formulae, pure-Python velocity-space relaxation tests with no compiled module or magnetic field, end-to-end tests of the full production tracer (guiding-centre orbits + Milstein noise), and cross-code validation against ASCOT5 [3] in both a controlled uniform plasma. Every number in Sections 9.2 and 9.3 was re-measured against the current tree and is reproducible from `tests/field/test_collisions.py` and `tests/field/test_collisions_local.py`. The cross-code comparison of Section 9.5 is the exception and is marked as such: it came from a driver script that has been removed pending a reworked validation example, so its numbers predate the current integrator and have not been regenerated.

9.1 Analytic Coefficient Unit Tests

`TestCollisionCoefficients` checks the Chandrasekhar function $G(x)$ and its derivative against known limits ($G(0) = 0$, $G(x \rightarrow \infty) \rightarrow 1/2x^2$, sign change of G' near $x \approx 0.968$), verifies $K < 0$ (drag always decelerates) and $\nu_D > 0$ for arbitrary finite backgrounds, and checks the expected scalings $K \propto n_b$, $K \propto q_a^2$, $\nu_D \propto n_b$, and additivity across species to $\lesssim 10\%$ (the residual comes from the $\ln \Lambda$ dependence on λ_D). `TestUnphysicalCoulombLog` confirms that $\ln \Lambda \leq 0$ raises an exception rather than returning silently unphysical coefficients (Section 4.2). These are deterministic, sub-second checks with no random component.

9.2 Local Velocity-Space Relaxation Tests

`tests/field/test_collisions_local.py` mirrors `collisions.h` in pure Python (no compiled module, no magnetic field) and integrates only the velocity-space SDE (Eqs. (26)–(27) with the orbital drifts set to zero) for an ensemble of $N = 4000$ particles. All collision rates scale linearly with the background density, so boosting $n_b \rightarrow 10^{23} \text{ m}^{-3}$ compresses the relaxation time to microseconds, letting the full suite run in about 10 seconds on a laptop. These tests are the most direct regression check on the Einstein-relation drag (Eq. (16)): an earlier, incorrect form of Q passed all analytic unit tests but drove a proton ensemble to $\langle E \rangle = 3.7 T_b$ instead of $1.5 T_b$ here, with a Kolmogorov–Smirnov (KS) p -value of 0 against the Maxwell speed law.

Test	$\langle E \rangle / T_b$ (target 1.5)	KS p -value	Pass
Cold start ($E_0 = 0.45 T_b$)	1.492	0.344	✓
Hot start ($E_0 = 4.5 T_b$)	1.537	0.251	✓
Stationarity (start at equilibrium)	1.445	0.064	✓

Table 2: Relaxation of a proton ensemble in a 1 keV proton background ($n_b = 10^{23} \text{ m}^{-3}$) to the Maxwellian, evolved for 8 relaxation times. KS p -values test the final speed distribution against the Maxwell speed law at the background temperature.

The same ensemble also verifies pitch isotropisation: starting from $\xi_0 = 1$, the relaxed ensemble gives $\langle \xi \rangle = 0.006$ (target 0) and $\langle \xi^2 \rangle = 0.336$ (target 1/3), consistent with convergence to $\xi \sim \mathcal{U}(-1, 1)$. A short-time drift check on 3.52 MeV α -particles against 10 keV electrons, over a window short enough that $K(v)$ is nearly constant (Doppler shift in speed $< 2\%$), gives $\frac{d\langle v \rangle}{dt} / K(v_0) = 0.984$ (target 1), confirming the deterministic drift term independently of the noise terms.

9.3 End-to-End Equilibration Tests (Production Tracer)

`TestMaxwellianEquilibration` and `TestPitchIsotropization` in `tests/field/test_collisions.py` exercise the full production path — `trace_particles_boozer_with_collisions`, i.e. the adaptive Dormand–Prince orbital integrator coupled to the Milstein noise step, compiled C++ coefficients, real (though simplified, `BoozerAnalytic`) magnetic geometry, and a `MaxToroidal-FluxStoppingCriterion` — rather than the idealised velocity-space-only integration of Section 9.2.

Test	Measured	Target	Pass
Hot start ($E_0 = 4.5 T_b$), $\langle E \rangle / T_b$	1.590	1.5	✓
Cold start ($E_0 = 0.45 T_b$), $\langle E \rangle / T_b$	1.799	1.5	✓
Pitch isotropisation, $\langle \xi \rangle$	+0.029	0	✓
Pitch isotropisation, $\langle \xi^2 \rangle$	0.331	1/3	✓
Pitch broadening ($\xi_0 = 0$), $\text{std}(\xi)$	0.418	> 0.05	✓

Table 3: End-to-end equilibration through the production tracer: 40 protons on a 1 keV proton background, $n_b = 10^{21} \text{ m}^{-3}$, 8 relaxation times, on the near-axis confining field of `TestMaxwellianEquilibration` ($R_0 = 8 \text{ m}$). All particles remain confined and finite throughout (no losses, no NaN states). The final row is measured separately by `TestPitchIsotropization`, on a different field and background. $N = 40$ gives a Monte Carlo spread of order 0.2 on $\langle E \rangle / T_b$, which is why these sit further from 3/2 than the $N = 4000$ velocity-space figures of Section 9.2.

The pitch-decay benchmark and its correction. A quantitative pitch-decay test against the naive expectation $\langle \xi(t) \rangle = \xi_0 e^{-\nu_D t}$ (the usual ASCOT5-style benchmark, which assumes v constant) failed the first time this test suite was run on Perlmutter: for 3.52 MeV α -particles on a single deuterium background, the drag rate is $(m_a/m_b) \nu_D = 2 \nu_D$, so the ensemble slows almost completely over the same window in which pitch decay is measured, and $\nu_D(v) \propto v^{-3}$ diverges as the particles thermalise — the ensemble fully isotropises ($\langle \xi \rangle \rightarrow 0$) rather than decaying exponentially at the fixed initial rate. The test was corrected to check the classical relation obtained by eliminating time between the drag and pitch equations,

$$\frac{\langle \xi(t_2) \rangle}{\langle \xi(t_1) \rangle} = \left(\frac{v(t_2)}{v(t_1)} \right)^{m_b/m_a}, \quad (37)$$

evaluated self-consistently against the simulation’s own mean speed (referencing the first saved snapshot rather than $t = 0$ cancels a geometric offset from sampling ξ along field-line-following orbits). A second, independent effect was found in the same debugging pass: in a toroidal (non-uniform- $|B|$) field the trapped-passing boundary can sit just below the injected pitch, so scattered particles fall into the trapped population where the orbit-averaged ξ is systematically pulled toward zero, roughly doubling the apparent decay rate relative to the uniform-plasma law — real neoclassical physics, but not what a velocity-space-operator test should measure. The test therefore uses a near-uniform- $|B|$ configuration ($\bar{\eta} \rightarrow 0$: no mirror force, no trapped population), giving

$$\frac{\langle \xi_2 \rangle / \langle \xi_1 \rangle}{(v_2/v_1)^{m_b/m_a}} = 1.038 \quad (\text{target } 1, \text{ tolerance } \pm 15\%),$$

with $v_2/v_1 = 0.498$ over the measurement window.

9.4 Slow Physics Suite (MPI, HPC)

Five tests marked `@pytest.mark.slow` run longer physical integrations ($N \sim 20$ –128 particles, 5–100 μs of physical time) that are impractical to run serially: an electron-drag signal-to-noise test (chosen specifically because ion-background drag is buried in noise for 3.52 MeV α -particles, whereas electron-background drag is enhanced by m_α/m_e), a cold-plasma energy-loss test, a

two-background superposition test, and the pitch-broadening/decay tests of Section 9.3. The tracer’s MPI support (particles distributed across ranks, results all-gathered identically on every rank, per-particle seeds rank-independent) lets this suite run on `tests/perlmutter/run_slow_tests.sh`: all five tests pass in 77s under 64 MPI ranks on Perlmutter, versus 3h45m running the same suite serially.

9.5 Cross-Code Validation Against ASCOT5: Uniform Plasma

This validation mirrors the alpha slowing-down branch of ASCOT5’s own built-in Coulomb-collision physics test: N fusion alphas (birth energy 3.52 MeV, isotropic pitch, shared random seed) slow down in a uniform 1 keV hydrogen plasma. The two codes use different magnetic geometries (ASCOT5’s analytical ITER-like circular tokamak; firm3d’s `BoozerAnalytic` at the same R_0, B_0 scale), but since the plasma is uniform the compared quantities — mean slowing-down time to 50 keV and $\langle E(t) \rangle$ — are geometry-insensitive, isolating the collision operator itself.

Provenance. Unlike every other number in this section, the figures below have *not* been regenerated against the current tree. They were produced by a driver script that has since been removed pending a reworked validation example, and they predate three changes that move them: the magnetic moment becoming a parameter of the orbit equations rather than a state variable, the collision drift moving into a sub-cycled kick (Section 6.4), and the correction of `ALPHA_PARTICLE_MASS` from the sum of free constituents to the measured value, which lowers the alpha mass by 0.76% and raises the birth speed by 0.38%. Treat the agreement recorded here as evidence about the coefficient set — which is unchanged and is pinned directly against ASCOT5 in Section 4 — rather than about the current integrator.

	firm3d	ASCOT5	Agreement
Mean slowing-down time [ms]	24.915 ± 0.027	24.885 ± 0.029	+0.1%, 0.7σ

Table 4: $N = 1024$ alphas, $n = 10^{20} \text{ m}^{-3}$ (the published ASCOT5 test-point density), $T = 1 \text{ keV}$. Also: $\max |\langle E(t) \rangle_{\text{firm3d}} - \langle E(t) \rangle_{\text{ASCOT5}}| / E_0 = 0.4\%$ over the full slowing-down window.

10 Validity and Limitations

- **Vacuum field:** electric fields, $E \times B$ drifts, and electromagnetic perturbations are not included in the collision model (though perturbations can be combined via the existing SAW tracing infrastructure).
- **Maxwellian backgrounds:** the collision operator assumes each background species has a local Maxwellian distribution function. Distorted distributions (e.g. beam-driven) are not supported.
- **Test-particle limit:** the EP density is assumed negligibly small compared to the background, so no back-reaction on the background is included.
- **Guiding-centre approximation:** valid when the Larmor radius $r_L = mv_\perp / (|q|B)$ is small compared to the scale lengths of B and the density/temperature profiles. For




Figure 1: firm3d vs. ASCOT5, uniform 1 keV plasma, $N = 1024$: mean energy $\langle E(t) \rangle$ against the analytic electron-drag estimate $v = v_0 e^{-t/\tau_{se}}$ (left); mean pitch $\langle \xi(t) \rangle$, shown for qualitative comparison only since it is geometry-sensitive (centre); and the distribution of per-particle slowing-down times to 50 keV against the analytic Spitzer estimate (right).

3.52 MeV α -particles in a 5 T stellarator $r_L \sim 5$ cm, comparable to device scale, so finite-Larmor-radius corrections may be non-negligible.

- **Zeroth-order GC collision operator:** the implementation follows Hirvijoki et al. [1] at zeroth order in ρ_*/L (Larmor radius over gradient length). The spatial diffusion term $D_X \propto D_\perp/\Omega_g^2$ present in the full operator [2] is neglected.
- **Coulomb logarithm held fixed in $\partial D_\parallel/\partial v$:** Eq. (15) treats $\ln \Lambda_{ab}$ as constant under differentiation, while the $\ln \Lambda_{ab}$ actually used varies with v through $\bar{v}^2$ (Eq. (9)). The spurious-drift terms in K (Eq. (17)) are then not quite the ones belonging to the $D_\parallel(v)$ that drives the noise, and the stationary solution of the speed SDE becomes the Maxwellian divided by $\ln \Lambda_{ab}(v)$ rather than the Maxwellian itself. Integrating the code's own K and $D_\parallel$ for a proton in 1 keV hydrogen at $n = 10^{21} \text{ m}^{-3}$ gives $\langle E \rangle/T_b = 1.469$ against the exact $3/2$, a -2.1% bias; the closed form above reproduces that number exactly. This is inherited, not accidental: ASCOT5's `mccc_coefs_clog` is likewise velocity-dependent and re-evaluated every step, while `mccc_coefs_dDpara` likewise carries c_{ab} (which contains $\ln \Lambda$) through the derivative, so the two codes agree coefficient for coefficient and share this bias. Adding the $\partial \ln \Lambda_{ab}/\partial v$ term would make the equilibrium exactly Maxwellian at the cost of departing from ASCOT5. Note the effect is on the thermal equilibrium, not on slowing-down: the uniform-plasma benchmark of Section 9.5 agrees with ASCOT5 to 0.7σ , as it should when both codes carry the same approximation.
- **Strong convergence of the noise step:** the diffusion matrix is not commutative, so omitting the Lévy-area terms leaves the scheme strong order $1/2$ rather than 1 in the v - ξ coupling; weak order 1 is retained. See Section 5.2.
- **Slow tests: `@pytest.mark.slow`** — five physics-validation tests require $N \sim 20$ particles integrated for 5–100 μs and are intended for HPC resources (Perlmutter or similar).

References

- [1] E. Hirvijoki, A. Brizard, A. Snicker, and T. Kurki-Suonio, “Monte Carlo implementation of a guiding-center Fokker-Planck kinetic equation,” *Phys. Plasmas* **20**, 092505 (2013). doi:10.1063/1.4820951
- [2] E. Hirvijoki et al., “ASCOT: Solving the kinetic equation of minority particle species in tokamak plasmas,” *Comput. Phys. Commun.* **185**, 1310 (2014). doi:10.1016/j.cpc.2013.10.019
- [3] ASCOT5 source code, `src/simulate/mccc/mccc_coefs.h`, function `mccc_coefs_clog`, https://github.com/ascot4fusion/ascot5/blob/main/src/simulate/mccc/mccc_coefs.h.
- [4] A. H. Boozer and G. Kuo-Petravic, “Monte Carlo evaluation of transport coefficients,” *Phys. Fluids* **24**, 851 (1981). doi:10.1063/1.863445
- [5] L. Spitzer, *Physics of Fully Ionized Gases*, 2nd ed. (Interscience, New York, 1962), Chapter 5.
- [6] J. D. Huba, *NRL Plasma Formulary* (Naval Research Laboratory, Washington DC, 2019). Available at <https://www.nrl.navy.mil/ppd/content/nrl-plasma-formulary>
- [7] S. Ichimaru, *Basic Principles of Plasma Physics: A Statistical Approach* (Benjamin, Reading MA, 1973).
- [8] S. Chandrasekhar, “Dynamical friction I: General considerations,” *Astrophys. J.* **97**, 255 (1943). doi:10.1086/144517
- [9] M. N. Rosenbluth, W. M. MacDonald, and D. L. Judd, “Fokker-Planck equation for an inverse-square force,” *Phys. Rev.* **107**, 1 (1957). doi:10.1103/PhysRev.107.1
- [10] B. A. Trubnikov, “Particle interactions in a fully ionized plasma,” in *Reviews of Plasma Physics*, Vol. 1, M. A. Leontovich, ed. (Consultants Bureau, New York, 1965), p. 105.
- [11] P. Helander and D. J. Sigmar, *Collisional Transport in Magnetized Plasmas* (Cambridge University Press, Cambridge, 2002), Chapter 3.
- [12] P. E. Kloeden and E. Platen, *Numerical Solution of Stochastic Differential Equations* (Springer, Berlin, 1992).
- [13] J. R. Dormand and P. J. Prince, “A family of embedded Runge-Kutta formulae,” *J. Comput. Appl. Math.* **6**, 19 (1980). doi:10.1016/0771-050X(80)90013-3